

Markov Decision Processes (MDPs)

Machine Learning – 10701/15781


Carlos Guestrin

Carnegie Mellon University

April 30th, 2007


©2005-2007 Carlos Guestrin

Thus far this semester



- Regression:
- Classification:
- Density estimation:

Learning to act




[Ng et al. '05]


- Reinforcement learning
- An agent
 - Makes sensor observations
 - Must select action
 - Receives rewards
 - positive for “good” states
 - negative for “bad” states

Learning to play backgammon

[Tesauro '95]

- Combines reinforcement learning with neural networks
- Played 300,000 games against itself
- Achieved grandmaster level!






Roadmap to learning about reinforcement learning

- When we learned about Bayes nets:
 - First talked about formal framework:
 - representation
 - inference
 - Then learning for BNs
- For reinforcement learning:
 - Formal framework
 - Markov decision processes
 - Then learning

FreeCraft



building


peasant

footman

Real-time Strategy Game

Peasants collect resources and build
Footmen attack enemies
Buildings train peasants and footmen

States and actions



- State space:

- Joint state \mathbf{x} of entire system

- Action space:

- Joint action $\mathbf{a} = \{a_1, \dots, a_n\}$ for all agents



States change over time

- Like an HMM, state changes over time
- Next state depends on current state and action selected
 - e.g., action="build castle" likely to lead to a state where you have a castle
- Transition model:
 - Dynamics of the entire system $P(x'|x,a)$



Some states and actions are better than others

- Each state x is associated with a reward
 - positive reward for successful attack
 - negative for loss
- Reward function:
 - Total reward $R(x)$




Discounted Rewards

An assistant professor gets paid, say, 20K per year.

How much, in total, will the A.P. earn in their life?

$$20 + 20 + 20 + 20 + 20 + \dots = \text{Infinity}$$



What's wrong with this argument?

Discounted Rewards



“A reward (payment) in the future is not worth quite as much as a reward now.”

- Because of chance of obliteration
- Because of inflation

Example:

Being promised \$10,000 next year is worth only 90% as much as receiving \$10,000 right now.

Assuming payment n years in future is worth only $(0.9)^n$ of payment now, what is the AP's **Future Discounted Sum of Rewards** ?

Discount Factors




People in economics and probabilistic decision-making do this all the time.

The “Discounted sum of future rewards” using discount factor γ is

$$\begin{aligned} & (\text{reward now}) + \\ & \gamma (\text{reward in 1 time step}) + \\ & \gamma^2 (\text{reward in 2 time steps}) + \\ & \gamma^3 (\text{reward in 3 time steps}) + \\ & \vdots \\ & \quad : \quad (\text{infinite sum}) \end{aligned}$$

The Academic Life

Assume Discount Factor $\gamma = 0.9$



Define:

V_A = Expected discounted future rewards starting in state A

V_B = Expected discounted future rewards starting in state B


V_T = “ “ “ “ “ “ “ “ T

V_S = “ “ “ “ “ “ “ “ S

V_D = “ “ “ “ “ “ “ “ D

How do we compute V_A , V_B , V_T , V_S , V_D ?

Computing the Future Rewards of an Academic



Assume Discount Factor $\gamma = 0.9$

Joint Decision Space

Markov Decision Process (MDP) Representation:

- State space:
 - Joint state \mathbf{x} of entire system
- Action space:
 - Joint action $\mathbf{a} = \{a_1, \dots, a_n\}$ for all agents
- Reward function:
 - Total reward $R(\mathbf{x}, \mathbf{a})$
 - sometimes reward can depend on action
- Transition model:
 - Dynamics of the entire system $P(\mathbf{x}'|\mathbf{x}, \mathbf{a})$



Policy

Policy: $\pi(\mathbf{x}) = \mathbf{a}$

At state \mathbf{x} ,
action \mathbf{a} for all
agents



$\pi(\mathbf{x}_0)$ = both peasants get wood



$\pi(\mathbf{x}_1)$ = one peasant builds
barrack, other gets gold




$\pi(\mathbf{x}_2)$ = peasants get gold,
footmen attack

Value of Policy

Value: $V_\pi(x)$

Expected long-term reward starting from x



Computing the value of a policy


$$V_{\pi}(x_0) = E_{\pi}[R(x_0) + \gamma R(x_1) + \gamma^2 R(x_2) + \gamma^3 R(x_3) + \gamma^4 R(x_4) + L]$$

- Discounted value of a state:

- value of starting from x_0 and continuing with policy π from then on

$$\begin{aligned} V_{\pi}(x_0) &= E_{\pi}[R(x_0) + \gamma R(x_1) + \gamma^2 R(x_2) + \gamma^3 R(x_3) + \dots] \\ &= E_{\pi}\left[\sum_{t=0}^{\infty} \gamma^t R(x_t)\right] \end{aligned}$$


- A recursion!

Computing the value of a policy 1 – the matrix inversion approach


$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$

- Solve by simple matrix inversion:


Computing the value of a policy 2 – iteratively


$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$


- If you have 1000,000 states, inverting a 1000,000x1000,000 matrix is hard!
- Can solve using a simple convergent iterative approach:
(a.k.a. dynamic programming)
 - Start with some guess V_0
 - Iteratively say:
 - $V_{t+1} = R + \gamma P_\pi V_t$
 - Stop when $\|V_{t+1} - V_t\|_1 \cdot \epsilon$
 - means that $\|V_\pi - V_{t+1}\|_1 \cdot \epsilon / (1 - \gamma)$

But we want to learn a Policy

- So far, told you how good a policy is...
- But how can we choose the best policy???
- Suppose there was only one time step:
 - world is about to end!!!
 - select action that maximizes reward!



Another recursion!



- Two time steps: address tradeoff
 - good reward now
 - better reward in the future

Unrolling the recursion

- Choose actions that lead to best value in the long run
 - Optimal value policy achieves optimal value V^*

$$V^*(x_0) = \max_{a_0} R(x_0, a_0) + \gamma E_{a_0} [\max_{a_1} R(x_1, a_1) + \gamma^2 E_{a_1} [\max_{a_2} R(x_2, a_2) + \dots]]$$

Bellman equation [Bellman 60]

- Evaluating policy π :


$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$

- Computing the optimal value V^* - Bellman equation

$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Optimal Long-term Plan

Optimal value
function $V^*(\mathbf{x})$



Optimal Policy: $\pi^*(\mathbf{x})$

$$Q^*(\mathbf{x}, \mathbf{a}) = R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Optimal policy:

$$\pi^*(\mathbf{x}) = \arg \max_{\mathbf{a}} Q^*(\mathbf{x}, \mathbf{a})$$

Interesting fact – Unique value

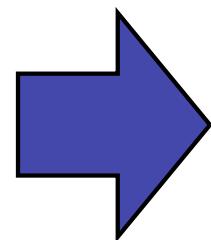

$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

- *Slightly surprising fact:* There is only one V^* that solves Bellman equation!
 - there may be many optimal policies that achieve V^*
- *Surprising fact:* optimal policies are good everywhere!!!


$$V_{\pi^*}(x) \geq V_\pi(x), \quad \forall x, \quad \forall \pi$$

Solving an MDP

Solve
Bellman
equation



Optimal
value $V^*(\mathbf{x})$



Optimal
policy $\pi^*(\mathbf{x})$


$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Bellman equation is non-linear!!!

Many algorithms solve the Bellman equations:

- Policy iteration [Howard '60, Bellman '57]
- Value iteration [Bellman '57]
- Linear programming [Manne '60]
- ...

Value iteration (a.k.a. dynamic programming) – the simplest of all


$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

- Start with some guess V_0
- Iteratively say:


- $V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$

- Stop when $\|V_{t+1} - V_t\|_1 \leq \varepsilon$
 - means that $\|V^* - V_{t+1}\|_1 \leq \varepsilon/(1-\gamma)$


A simple example

You run a startup company.

In every state you must choose between Saving money or Advertising.





Let's compute $V_t(x)$ for our example



$$V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

Let's compute $V_t(x)$ for our example



t	$V_t(\text{PU})$	$V_t(\text{PF})$	$V_t(\text{RU})$	$V_t(\text{RF})$
1	0	0	10	10
2	0	4.5	14.5	19
3	2.03	6.53	25.08	18.55
4	3.852	12.20	29.63	19.26
5	7.22	15.07	32.00	20.40
6	10.03	17.65	33.58	22.43

$$V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

Policy iteration – Another approach for computing π^*

- Start with some guess for a policy π_0
- Iteratively say:

- evaluate policy:

$$V_t(\mathbf{x}) = R(\mathbf{x}, \mathbf{a} = \pi_t(\mathbf{x})) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a} = \pi_t(\mathbf{x})) V_t(\mathbf{x}')$$

- improve policy:

$$\pi_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

- Stop when
 - policy stops changing
 - usually happens in about 10 iterations
 - or $\|V_{t+1} - V_t\|_1 \cdot \epsilon$
 - means that $\|V^* - V_{t+1}\|_1 \cdot \epsilon / (1 - \gamma)$

Policy Iteration & Value Iteration: Which is best ???

It depends.

Lots of actions? Choose **Policy Iteration**

Already got a fair policy? **Policy Iteration**

Few actions, acyclic? **Value Iteration**

Best of Both Worlds:

Modified Policy Iteration [Puterman]

...a simple mix of value iteration and policy iteration

3rd Approach

Linear Programming

LP Solution to MDP

[Manne '60]

Value computed by linear programming:

$$\text{minimize: } \sum_{\mathbf{x}} V(\mathbf{x})$$

$$\text{subject to : } \begin{cases} V(\mathbf{x}) \geq R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V(\mathbf{x}') \\ \forall \mathbf{x}, \mathbf{a} \end{cases}$$


- One variable $V(\mathbf{x})$ for each state
- One constraint for each state \mathbf{x} and action \mathbf{a}
- Polynomial time solution

What you need to know

- What's a Markov decision process
 - state, actions, transitions, rewards
 - a policy
 - value function for a policy
 - computing V_π
- Optimal value function and optimal policy
 - Bellman equation
- Solving Bellman equation
 - with value iteration, policy iteration and linear programming

Acknowledgment

- This lecture contains some material from Andrew Moore's excellent collection of ML tutorials:
 - <http://www.cs.cmu.edu/~awm/tutorials>



Reinforcement Learning

Machine Learning – 10701/15781
Carlos Guestrin
Carnegie Mellon University

April 30th, 2007

©2005-2007 Carlos Guestrin

The Reinforcement Learning task



World: You are in state 34.

Your immediate reward is 3. You have possible 3 actions.

Robot: I'll take action 2.

World: You are in state 77.

Your immediate reward is -7. You have possible 2 actions.

Robot: I'll take action 1.


World: You're in state 34 (again).

Your immediate reward is 3. You have possible 3 actions.

Formalizing the (online) reinforcement learning problem

- Given a set of states X and actions A
 - in some versions of the problem size of X and A unknown
- Interact with world at each time step t :
 - world gives state x_t and reward r_t
 - you give next action a_t
- **Goal:** (quickly) learn policy that (approximately) maximizes long-term expected discounted reward

The “Credit Assignment” Problem



I'm in state 43, reward = 0, action = 2

“ “ “ 39, “ = 0, “ = 4

“ “ “ 22, “ = 0, “ = 1

“ “ “ 21, “ = 0, “ = 1

“ “ “ 21, “ = 0, “ = 1

“ “ “ 13, “ = 0, “ = 2


“ “ “ 54, “ = 0, “ = 2

“ “ “ 26, “ = 100,

Yippee! I got to a state with a big reward! But which of my actions along the way actually helped me get there??

This is the **Credit Assignment** problem.

Exploration-Exploitation tradeoff



- You have visited part of the state space and found a reward of 100
 - is this the best I can hope for???
- **Exploitation:** should I stick with what I know and find a good policy w.r.t. this knowledge?
 - at the risk of missing out on some large reward somewhere
- **Exploration:** should I look for a region with more reward?
 - at the risk of wasting my time or collecting a lot of negative reward

Two main reinforcement learning approaches

- Model-based approaches:

- explore environment ! learn model ($P(x'|x,a)$ and $R(x,a)$) (almost) everywhere
 - use model to plan policy, MDP-style
 - approach leads to strongest theoretical results
 - works quite well in practice when state space is manageable

- Model-free approach:

- don't learn a model ! learn value function or policy directly
 - leads to weaker theoretical results
 - often works well when state space is large