

# Neural Networks

Machine Learning – 10701/15781

Carlos Guestrin, XOXO

Carnegie Mellon University

February 14<sup>th</sup>, 2007

©2005-2007 Carlos Guestrin

1

## Logistic regression

- P(Y|X) represented by:

$$\begin{aligned} P(Y = 1 | \mathbf{x}, W) &= \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}} \\ &= g(w_0 + \sum_i w_i x_i) \end{aligned}$$

- Learning rule – MLE:

Compute derivative

$$\frac{\partial \ell(W)}{\partial w_i} = \sum_j x_i^j [y^j - P(Y^j = 1 | \mathbf{x}^j, W)]$$


$$= \sum_j x_i^j [y^j - g(w_0 + \sum_i w_i x_i^j)]$$

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \sum_j x_i^j \delta^j$$

$$\delta^j = y^j - g(w_0 + \sum_i w_i x_i^j)$$

©2005-2007 Carlos Guestrin


2




# Sigmoid

$$g(w_0 + \sum_i w_i x_i) = \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}}$$


$w_0=2, w_1=1$



$w_0=0, w_1=1$



$w_0=0, w_1=0.5$



©2005-2007 Carlos Guestrin

3

# Perceptron as a graph


input node



$$g(w_0 + \sum_i w_i x_i) = \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}}$$

output


$$y = g(w_0 + \sum_i w_i x_i)$$



©2005-2007 Carlos Guestrin

4

## Linear perceptron classification region



$$g(w_0 + \sum_i w_i x_i) = \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}}$$

$w_0 + \sum_i w_i x_i > 0$

$\Theta$

$w_0 + \sum_i w_i x_i < 0$

©2005-2007 Carlos Guestrin

5

## Optimizing the perceptron

$$\begin{aligned} \frac{\partial f(g(\mathbf{w}))}{\partial w_j} &= \\ &= \frac{\partial f(g(\mathbf{w}))}{\partial g} \cdot \frac{\partial g}{\partial w_j} \end{aligned}$$

- Trained to minimize sum-squared error

$$\ell(\mathbf{w}) = \frac{1}{2} \sum_{j=1}^m [y_j - g(w_0 + \sum_i w_i x_i^j)]^2$$

truth prediction

$$\frac{\partial \ell(\mathbf{w})}{\partial w_m} = \frac{\partial}{\partial w_m} \left[ \frac{1}{2} \sum_j [y_j - g(w_0 + \sum_i w_i x_i^j)]^2 \right]$$

$$\begin{aligned} &= \frac{1}{2} \sum_j \frac{\partial}{\partial w_m} [y_j - g(w_0 + \sum_i w_i x_i^j)]^2 \end{aligned}$$

©2005-2007 Carlos Guestrin

6

## Derivative of sigmoid

$$\frac{\partial \ell(W)}{\partial w_\mu} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j g'(w_0 + \sum_i w_i x_i^j)$$

*diff between truth and prediction*

*magnitude of derivative*  
*derivative of logistic function*

$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g'(x) = \frac{\partial g(x)}{\partial x} = g(x)(1 - g(x))$$

*derivative large near decision boundary*

$$\frac{\partial \ell}{\partial w_\mu} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j g'(w_0 + \sum_i w_i x_i^j)$$

©2005-2007 Carlos Guestrin

7

## The perceptron learning rule

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \sum_j x_i^j \delta^j$$

*learning rate*  
*strength fraction*

*pushes me towards right answer*  
*push away from boundary*

$$\delta^j = [y^j - g(w_0 + \sum_i w_i x_i^j)] g^j (1 - g^j)$$

$$g^j = g(w_0 + \sum_i w_i x_i^j)$$

*both unhappy  
when classification is wrong*

- Compare to MLE:

$$w_i \leftarrow w_i + \eta \sum_j x_i^j \delta^j$$

$$\delta^j = [y^j - g(w_0 + \sum_i w_i x_i^j)]$$

©2005-2007 Carlos Guestrin

8

## Perceptron, linear classification, Boolean functions

- Can learn  $x_1 \vee x_2 = y$

$$1 \xrightarrow{-0.5}$$

$$x_1 \xrightarrow{1} y$$

- Can learn  $x_1 \wedge x_2 = y$

$$\begin{array}{c} 1 \xrightarrow{-1.5} \\ x_1 \xrightarrow{1} y \\ x_2 \end{array}$$

- Can learn any conjunction or disjunction

$$x_1 \vee x_2 \vee \dots \vee x_n$$

$$1 \xrightarrow{-0.5}$$

$$x_1 \xrightarrow{1} y$$

$$\vdots \quad \vdots$$

$$x_1 \wedge x_2 \dots \wedge x_n$$

$$1 \xrightarrow{-1.5}$$

$$x_1 \xrightarrow{1} y$$

$$\vdots \quad \vdots$$

$$x_n \xrightarrow{1}$$

©2005-2007 Carlos Guestrin

9


## Perceptron, linear classification, Boolean functions

- Can learn majority

$$\sum_i x_i > \frac{n}{2}$$

- Can perceptrons do everything?

$$\text{XOR } \equiv y = x_1 \wedge \neg x_2 \vee \neg x_1 \wedge x_2$$




©2005-2007 Carlos Guestrin

10

## Going beyond linear classification

- Solving the XOR problem



©2005-2007 Carlos Guestrin

11


## Hidden layer

- Perceptron:  $\underline{out(\mathbf{x})} = g(w_0 + \sum_i w_i x_i)$

- 1-hidden layer:

$$out(\mathbf{x}) = g\left(w_0 + \sum_k w_k g\left(w_0^k + \sum_i w_i^k x_i\right)\right)$$


$w_k$



©2005-2007 Carlos Guestrin

12

## Example data for NN with hidden layer



A target function:

Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned??

13

## Learned weights for hidden layer

A network:



activation of  $V_1, V_2, V_3$   
correspond to  
binary encoding

Learned hidden layer representation:

Input	Hidden Values	Output
10000000	→ .89 .04 .08	→ 10000000
01000000	→ .01 .11 .88	→ 01000000
00100000	→ .01 .97 .27	→ 00100000
00010000	→ .99 .97 .71	→ 00010000
00001000	→ .03 .05 .02	→ 00001000
00000100	→ .22 .99 .99	→ 00000100
00000010	→ .80 .01 .98	→ 00000010
00000001	→ .60 .94 .01	→ 00000001

14

## NN for images

left strt right up output

inputs



Typical input images

90% accurate learning head pose, and recognizing 1-of-20 faces


15

## Weights in NN for images

left strt right up



Learned Weights



weights for hidden nodes  
white → positive  
gray → zero  
black → negative



Typical input images

©2005-2007 Carlos Guestrin


16

## Forward propagation for 1-hidden layer - Prediction

- 1-hidden layer:

$$\underline{out(x)} = g \left( w_0 + \sum_k w_k g(w_0^k + \sum_i w_i^k x_i) \right)$$

w's are given  
want to classify input x



©2005-2007 Carlos Guestrin

17

## Gradient descent for 1-hidden layer – Back-propagation: Computing $\frac{\partial \ell(W)}{\partial w_k}$

$$\ell(W) = \frac{1}{2} \sum_j [y^j - out(x^j)]^2$$

Dropped  $w_0$  to make derivation simpler

$$\underline{out(x)} = g \left( \sum_{k'} w_{k'} g \left( \sum_{i'} w_{i'}^k x_{i'} \right) \right)$$

$$\frac{\partial out}{\partial w_k} = \frac{\partial f}{\partial w_k} \cdot \frac{\partial g}{\partial f}$$

$$g' = g(1-g)$$

$$\frac{\partial \ell(W)}{\partial w_k} = \sum_{j=1}^m [y_j - out(x^j)] \frac{\partial out(x^j)}{\partial w_k}$$

$$\frac{\partial out(x^j)}{\partial w_k} = g \left( \sum_{i'} w_{i'}^k x_{i'} \right)$$

$$g' \left( \sum_{i'} w_{i'}^k g \left( \sum_{j'} w_{j'}^{i'} x_{j'} \right) \right)$$

$$= V_k out(x) (1-out(x))$$

©2005-2007 Carlos Guestrin

18

## Gradient descent for 1-hidden layer – Back-propagation: Computing $\frac{\partial \ell(W)}{\partial w_i^k}$

$$\ell(W) = \frac{1}{2} \sum_j [y^j - out(\mathbf{x}^j)]^2$$

Dropped  $w_0$  to make derivation simpler

$$out(\mathbf{x}) = g\left(\sum_{k'} w_{k'} g\left(\sum_{i'} w_{i'}^{k'} x_{i'}\right)\right)$$

$$\frac{\partial \ell(W)}{\partial w_i^k} = \sum_{j=1}^m -[y - out(\mathbf{x}^j)] \frac{\partial out(\mathbf{x}^j)}{\partial w_i^k}$$


$$\frac{\partial out(\mathbf{x})}{\partial w_i^k} = x_i \cdot g'(\sum_i w_i^k x_i) = g'(\sum_k w_k g(\sum_i w_{ii}^k x_{ii}))$$

$$= x_i \cdot v_k(1-v_k) \cdot out(x) \cdot (1-out(x))$$

©2005-2007 Carlos Guestrin

19

## Multilayer neural networks



©2005-2007 Carlos Guestrin

20

## Forward propagation – prediction

- Recursive algorithm
- Start from input layer
- Output of node  $V_k$  with parents  $U_1, U_2, \dots$ :

$$V_k = g\left(\sum_i w_i^k U_i\right)$$

©2005-2007 Carlos Guestrin

21

## Back-propagation – learning

- Just gradient descent!!!
- Recursive algorithm for computing gradient
- For each example
  - Perform forward propagation
  - Start from output layer
  - Compute gradient of node  $V_k$  with parents  $U_1, U_2, \dots$
  - Update weight  $w_i^k$

©2005-2007 Carlos Guestrin

22

## Many possible response functions

- Sigmoid  $g(w_0 + \sum_i w_i x_i)$



- Linear  $w_0 + \sum_i w_i x_i$




- Exponential  $e^{w_0 + \sum_i w_i x_i}$

$$e^{w_0 + \sum_i w_i x_i}$$



- Gaussian  $e^{-\frac{(w_0 + \sum_i w_i x_i)^2}{\sigma^2}}$

$$e^{-\frac{(w_0 + \sum_i w_i x_i)^2}{\sigma^2}}$$



- ... Threshold

:



©2005-2007 Carlos Guestrin

23

## Convergence of backprop

- Perceptron leads to convex optimization
  - Gradient descent reaches **global minima**



- Multilayer neural nets **not convex**
  - Gradient descent gets stuck in local minima
  - Hard to set learning rate
  - Selecting number of hidden units and layers = fuzzy process
  - NNs falling in disfavor in last few years
  - We'll see later in semester, kernel trick is a good alternative
  - Nonetheless, neural nets are one of the most used ML approaches




©2005-2007 Carlos Guestrin

24

## ~~Training set error~~ Overfitting

- Neural nets represent complex functions
  - Output becomes more complex with gradient steps



©2005-2007 Carlos Guestrin

25

## Overfitting

- Output fits training data "too well"
  - Poor test set accuracy
- Overfitting the training data
  - Related to bias-variance tradeoff
  - One of central problems of ML
- Avoiding overfitting?
  - More training data
  - Regularization, regularize w.r.t.  $w^2$  (typically)  
momentum, ..
  - Early stopping

©2005-2007 Carlos Guestrin

26

# What you need to know about neural networks

- Perceptron:
  - Representation
  - Perceptron learning rule
  - Derivation
- Multilayer neural nets
  - Representation
  - Derivation of backprop
  - Learning rule
- Overfitting
  - Definition
  - Training set versus test set
  - Learning curve

©2005-2007 Carlos Guestrin

27

## Instance-based Learning


Machine Learning – 10701/15781  
Carlos Guestrin  
Carnegie Mellon University

February 14<sup>th</sup>, 2007

©2005-2007 Carlos Guestrin

28


## Why not just use Linear Regression?



©2005-2007 Carlos Guestrin

29


## Using data to predict new data



©2005-2007 Carlos Guestrin

30

# Nearest neighbor



©2005-2007 Carlos Guestrin

31

## Univariate 1-Nearest Neighbor

Given datapoints  $(x_1, y_1), (x_2, y_2) \dots (x_N, y_N)$ , where we assume  $y_i = f(x_i)$  for some unknown function  $f$ .


Given query point  $x_q$ , your job is to predict  $\hat{y} \approx f(x_q)$   
Nearest Neighbor:

1. Find the closest  $x_i$  in our set of datapoints

$$i(nn) = \operatorname{argmin}_i |x_i - x_q|$$

2. Predict  $\hat{y} = y_{i(nn)}$

Here's a dataset with one input, one output and four datapoints.




©2005-2007 Carlos Guestrin

32

## 1-Nearest Neighbor is an example of.... Instance-based learning

A function approximator  
that has been around  
since about 1910.

To make a prediction,  
search database for  
similar datapoints, and fit  
with the local points.



### Four things make a memory based learner:

- A distance metric
- How many nearby neighbors to look at?
- A weighting function (optional)
- How to fit with the local points?

©2005-2007 Carlos Guestrin

33

## 1-Nearest Neighbor

### Four things make a memory based learner:

1. A distance metric  
**Euclidian (and many more)**
2. How many nearby neighbors to look at?  
**One**
3. A weighting function (optional)  
**Unused**
4. How to fit with the local points?  
**Just predict the same output as the nearest neighbor.**

©2005-2007 Carlos Guestrin

34

## Multivariate 1-NN examples

Regression

Classification

©2005-2007 Carlos Guestrin


35

## Multivariate distance metrics


Suppose the input vectors  $x_1, x_2, \dots, x_n$  are two dimensional:

$x_1 = (x_{11}, x_{12}), x_2 = (x_{21}, x_{22}), \dots, x_N = (x_{N1}, x_{N2})$ .

One can draw the nearest-neighbor regions in input space.



$$Dist(x_i, x_j) = (x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2$$



$$Dist(x_i, x_j) = (x_{i1} - x_{j1})^2 + (3x_{i2} - 3x_{j2})^2$$

The relative scalings in the distance metric affect region shapes.

©2005-2007 Carlos Guestrin

36