

Dimensionality reduction (cont.)

Machine Learning – 10701/15781

Carlos Guestrin

Carnegie Mellon University

April 25th, 2007

©2005-2007 Carlos Guestrin

Lower dimensional projections

- Rather than picking a subset of the features, we can new features that are combinations of existing features

x_1, x_2, x_3

$x_2^2, x_1 x_3 \dots$

Select
some of these

projection:

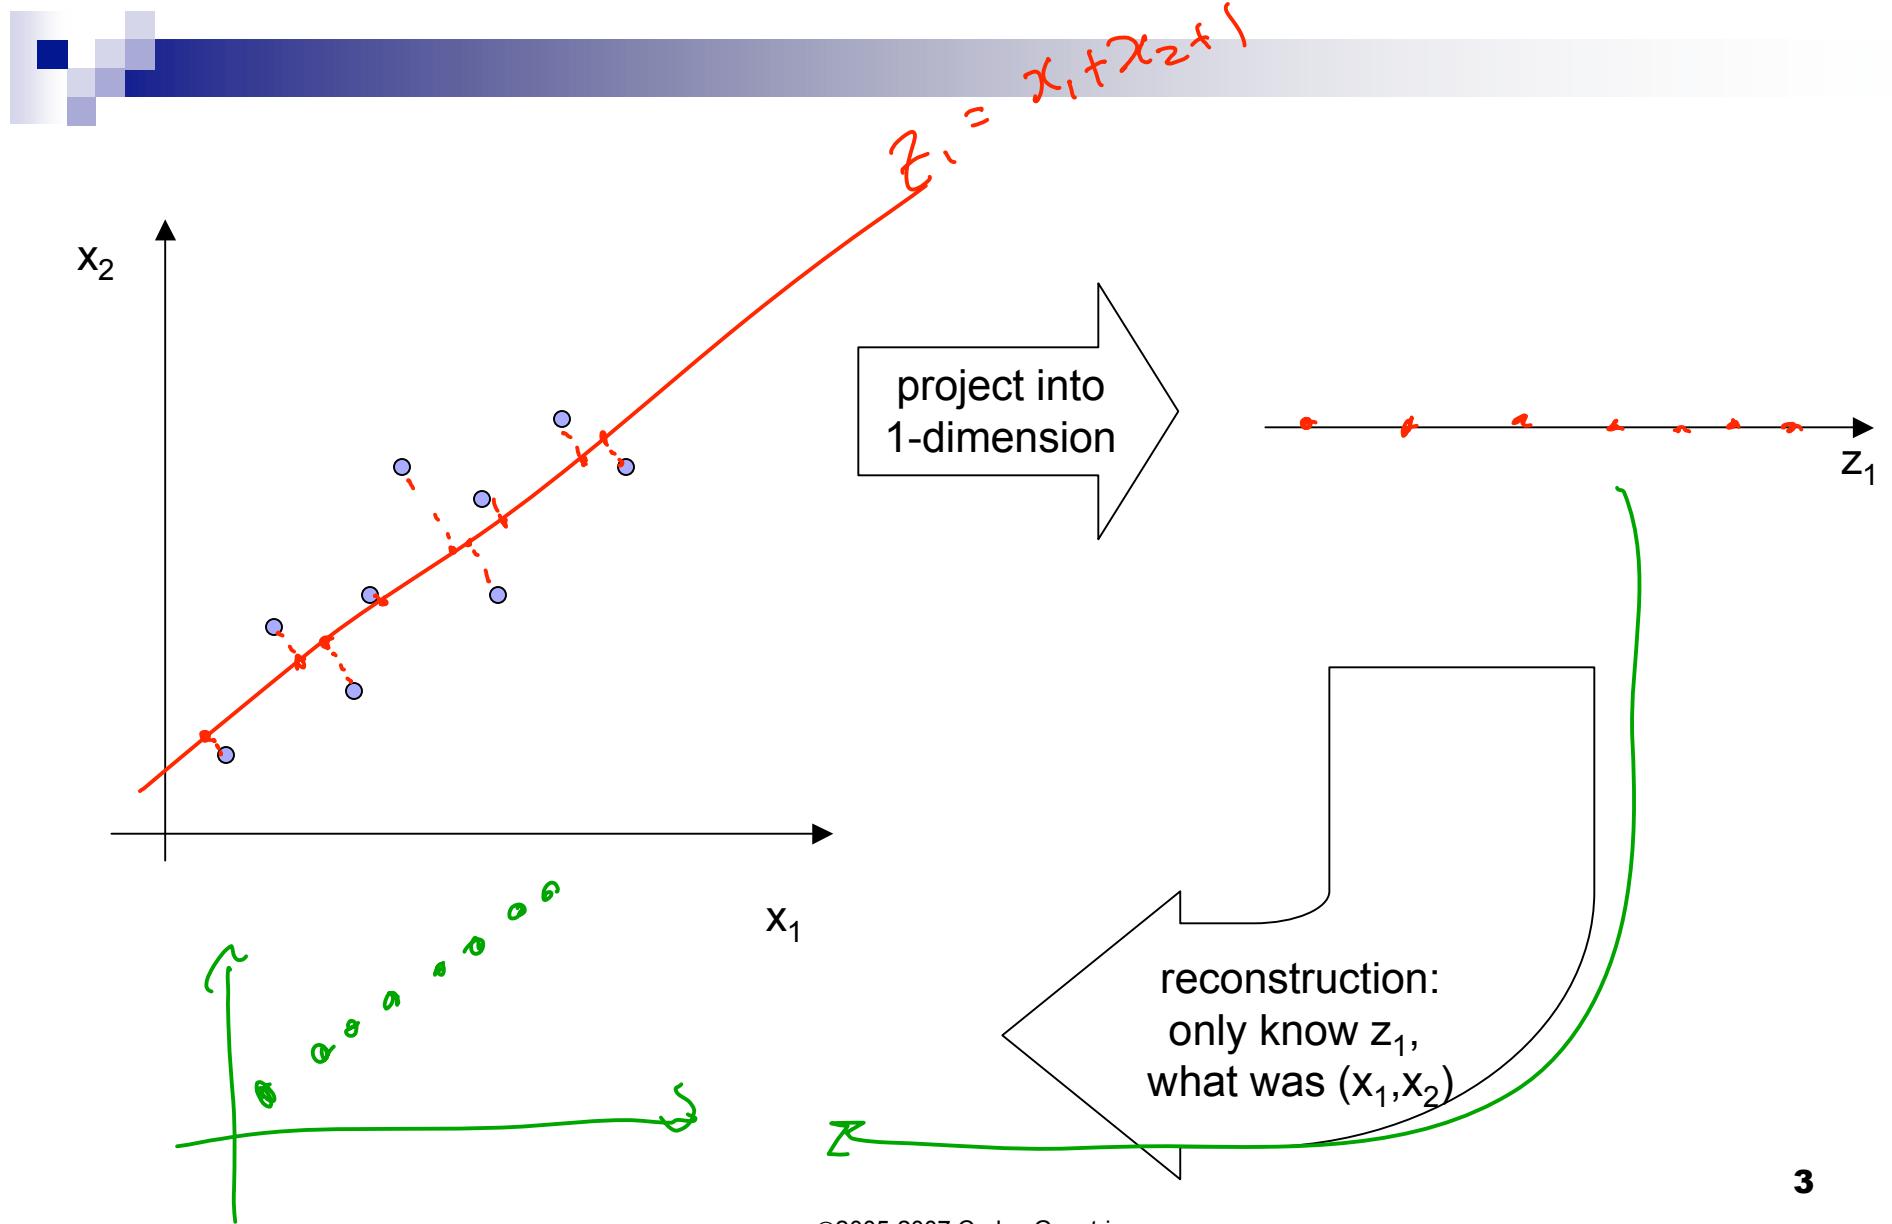
new feature, e.g.,

$$\underline{x_{\text{new}}} = 0.5x_1 - 0.75x_2 + 0.92x_3 \dots$$

↑
new
feature

- Let's see this in the unsupervised setting
 - just X , but no Y

Linear projection and reconstruction



Principal component analysis – basic idea

- Project n-dimensional data into k-dimensional space while preserving information:
 - e.g., project space of 10000 words into 3-dimensions
 - e.g., project 3-d into 2-d
- Choose projection with minimum reconstruction error

Linear projections, a review

- Project a point into a (lower dimensional) space:
 - **point:** $\mathbf{x} = (x_1, \dots, x_n)$
 - **select a basis** – set of basis vectors – $(\mathbf{u}_1, \dots, \mathbf{u}_k)$
 - we consider orthonormal basis:
 - $\mathbf{u}_i \cdot \mathbf{u}_i = 1$, and $\mathbf{u}_i \cdot \mathbf{u}_j = 0$ for $i \neq j$
 - **select a center** – $\bar{\mathbf{x}}$, defines offset of space
 - **best coordinates** in lower dimensional space defined by dot-products: (z_1, \dots, z_k) , $z_i = (\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{u}_i$
 - minimum squared error

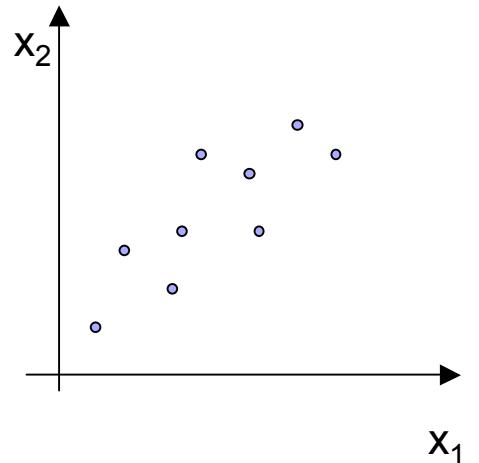
PCA finds projection that minimizes reconstruction error

- Given m data points: $\mathbf{x}^i = (x_1^i, \dots, x_n^i)$, $i=1\dots m$
- Will represent each point as a projection:

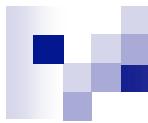
- $\hat{\mathbf{x}}^i = \bar{\mathbf{x}} + \sum_{j=1}^k z_j^i \mathbf{u}_j$ where: $\bar{\mathbf{x}} = \frac{1}{m} \sum_{i=1}^m \mathbf{x}^i$ and $z_j^i = (\mathbf{x}^i - \bar{\mathbf{x}}) \cdot \mathbf{u}_j$

- PCA:
 - Given $k \cdot n$, find $(\mathbf{u}_1, \dots, \mathbf{u}_k)$ minimizing reconstruction error:

$$error_k = \sum_{i=1}^m (\mathbf{x}^i - \hat{\mathbf{x}}^i)^2$$



Understanding the reconstruction error



- Note that \mathbf{x}^i can be represented exactly by n-dimensional projection:

$$\mathbf{x}^i = \bar{\mathbf{x}} + \sum_{j=1}^n z_j^i \mathbf{u}_j$$

- Rewriting error:

$$\hat{\mathbf{x}}^i = \bar{\mathbf{x}} + \sum_{j=1}^k z_j^i \mathbf{u}_j$$

$$z_j^i = (\mathbf{x}^i - \bar{\mathbf{x}}) \cdot \mathbf{u}_j$$

□ Given $k \cdot n$, find $(\mathbf{u}_1, \dots, \mathbf{u}_k)$ minimizing reconstruction error:

$$error_k = \sum_{i=1}^m (\mathbf{x}^i - \hat{\mathbf{x}}^i)^2$$

Reconstruction error and covariance matrix

$$error_k = \sum_{i=1}^m \sum_{j=k+1}^n [\mathbf{u}_j \cdot (\mathbf{x}^i - \bar{\mathbf{x}})]^2$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (\mathbf{x}^i - \bar{\mathbf{x}})(\mathbf{x}^i - \bar{\mathbf{x}})^T$$

Minimizing reconstruction error and eigen vectors

- Minimizing reconstruction error equivalent to picking orthonormal basis $(\mathbf{u}_1, \dots, \mathbf{u}_n)$ minimizing:

$$error_k = \sum_{j=k+1}^n \mathbf{u}_j^T \Sigma \mathbf{u}_j$$

- Eigen vector:

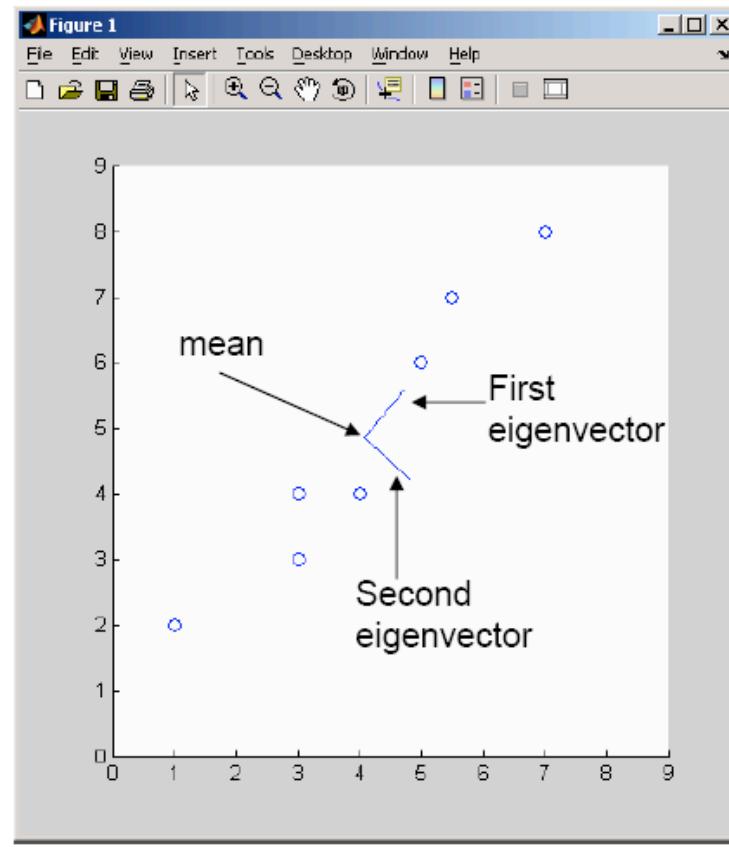
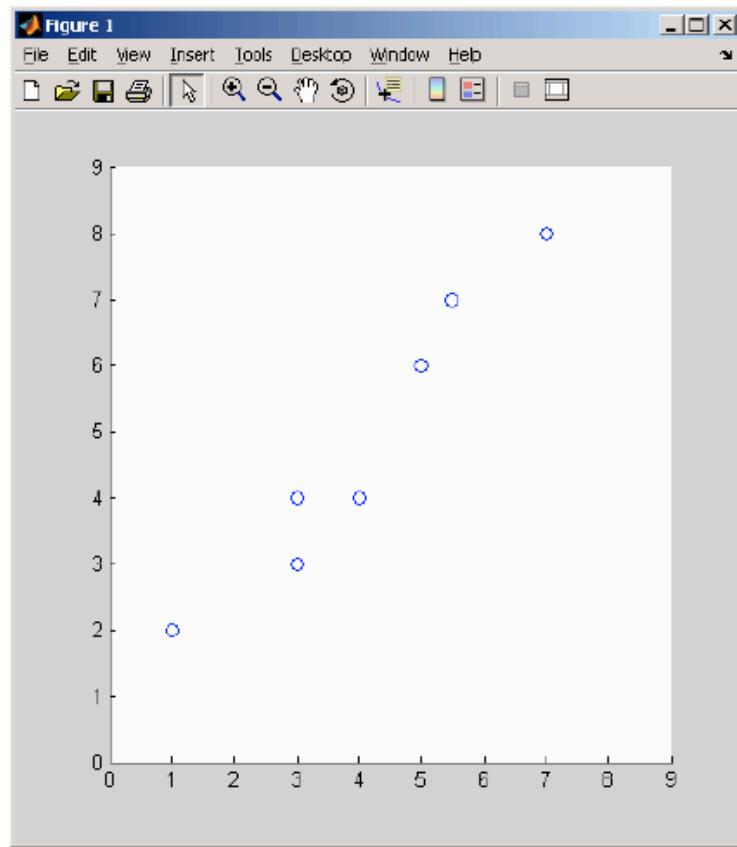
- Minimizing reconstruction error equivalent to picking $(\mathbf{u}_{k+1}, \dots, \mathbf{u}_n)$ to be eigen vectors with smallest eigen values

Basic PCA algorithm

- Start from m by n data matrix \mathbf{X}
- **Recenter:** subtract mean from each row of \mathbf{X}
 - $\mathbf{X}_c \leftarrow \mathbf{X} - \bar{\mathbf{X}}$
- **Compute covariance matrix:**
 - $\Sigma \leftarrow 1/m \mathbf{X}_c^T \mathbf{X}_c$
- Find **eigen vectors and values** of Σ
- **Principal components:** k eigen vectors with highest eigen values

PCA example

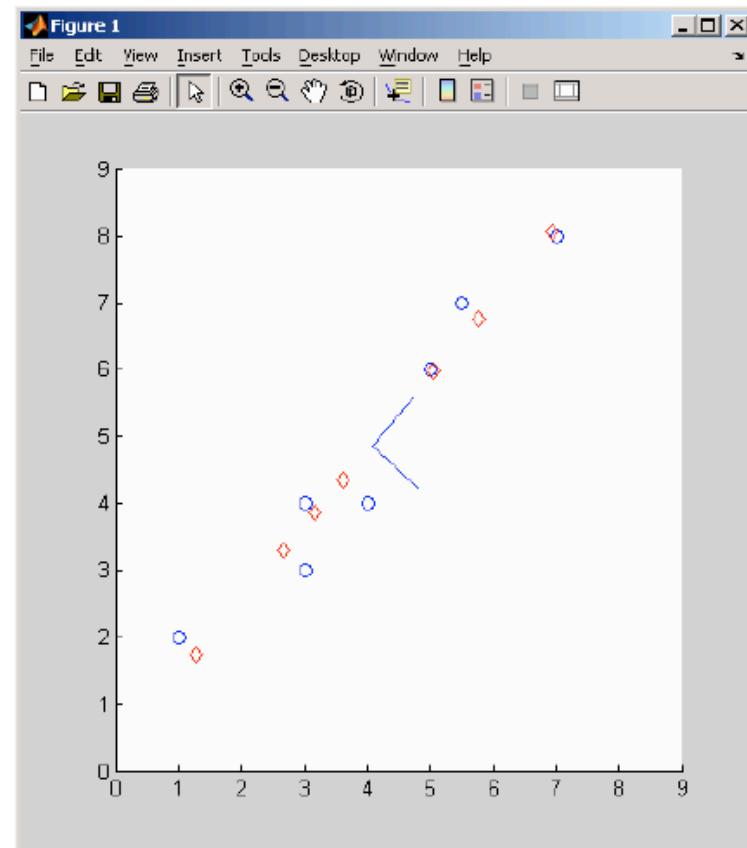
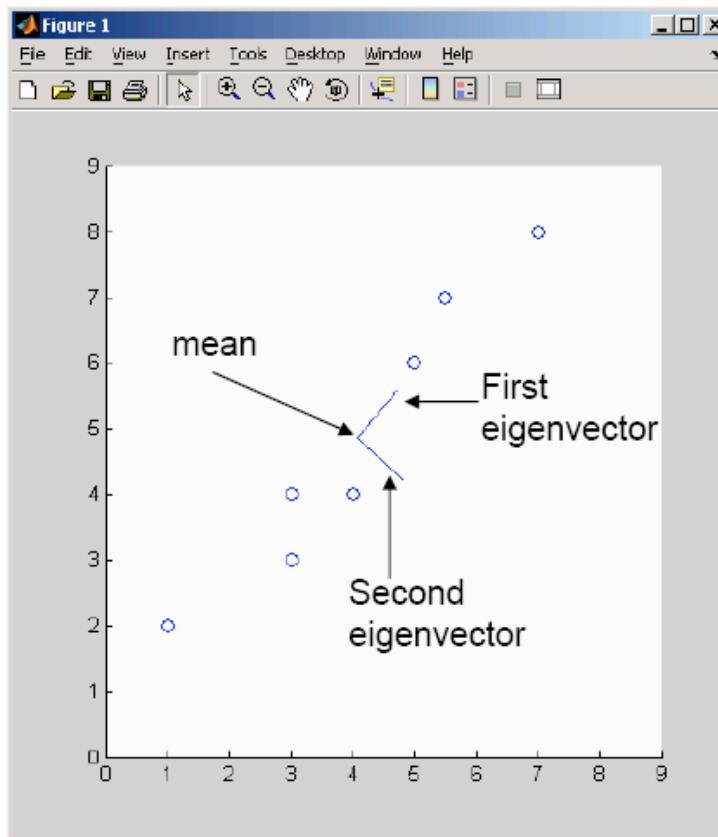
$$\hat{\mathbf{x}}^i = \bar{\mathbf{x}} + \sum_{j=1}^k z_j^i \mathbf{u}_j$$



PCA example – reconstruction

$$\hat{\mathbf{x}}^i = \bar{\mathbf{x}} + \sum_{j=1}^k z_j^i \mathbf{u}_j$$

only used first principal component



Eigenfaces [Turk, Pentland '91]

- Input images:



- Principal components:



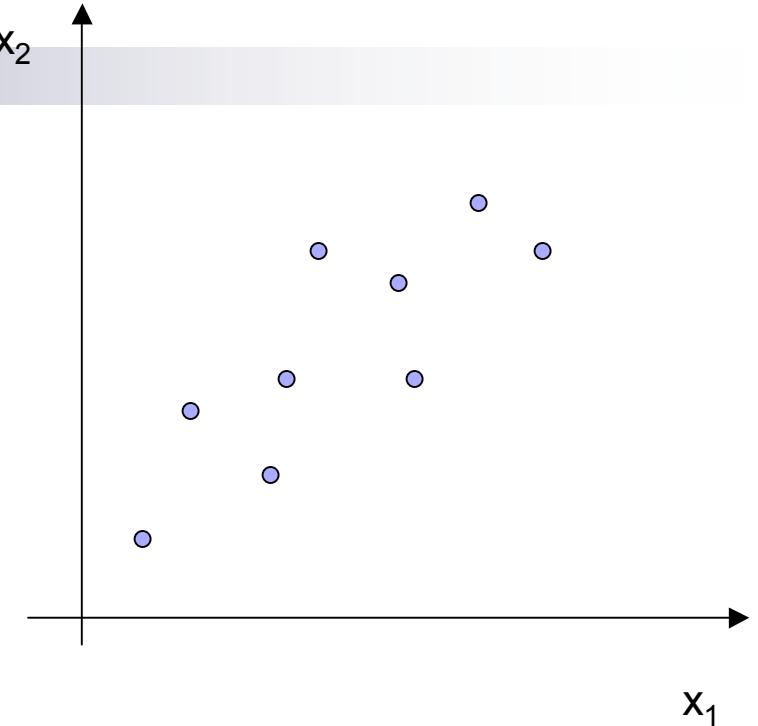
Eigenfaces reconstruction

- Each image corresponds to adding 8 principal components:



Relationship to Gaussians

- PCA assumes data is Gaussian
 - $\mathbf{x} \sim N(\bar{\mathbf{x}}; \Sigma)$
- Equivalent to weighted sum of simple Gaussians:
$$\mathbf{x} = \bar{\mathbf{x}} + \sum_{j=1}^n z_j \mathbf{u}_j; \quad z_j \sim N(0; \sigma_j^2)$$
- Selecting top k principal components equivalent to lower dimensional Gaussian approximation:
$$\mathbf{x} \approx \bar{\mathbf{x}} + \sum_{j=1}^k z_j \mathbf{u}_j + \varepsilon; \quad z_j \sim N(0; \sigma_j^2)$$
 - $\varepsilon \sim N(0; \sigma^2)$, where σ^2 is defined by error _{k}



Scaling up

- Covariance matrix can be really big!
 - Σ is n by n
 - 10000 features ! $|\Sigma|$
 - finding eigenvectors is very slow...
- Use singular value decomposition (SVD)
 - finds top k eigenvectors
 - great implementations available, e.g., Matlab svd

SVD

- Write $\mathbf{X} = \mathbf{W} \mathbf{S} \mathbf{V}^T$
 - $\mathbf{X} \leftarrow$ data matrix, one row per datapoint
 - $\mathbf{W} \leftarrow$ weight matrix, one row per datapoint – coordinate of \mathbf{x}^i in eigenspace
 - $\mathbf{S} \leftarrow$ singular value matrix, diagonal matrix
 - in our setting each entry is eigenvalue λ_j
 - $\mathbf{V}^T \leftarrow$ singular vector matrix
 - in our setting each row is eigenvector \mathbf{v}_j

PCA using SVD algorithm

- Start from m by n data matrix \mathbf{X}
- **Recenter:** subtract mean from each row of \mathbf{X}
 - $\mathbf{X}_c \leftarrow \mathbf{X} - \bar{\mathbf{x}}$
- Call SVD algorithm on \mathbf{X}_c – ask for k singular vectors
- **Principal components:** k singular vectors with highest singular values (rows of \mathbf{V}^T)
 - **Coefficients** become:



Using PCA for dimensionality reduction in classification

- Want to learn $f: \mathbf{X} \rightarrow \mathbf{Y}$
 - $\mathbf{X} = \langle X_1, \dots, X_n \rangle$
 - but some features are more important than others
- **Approach:** Use PCA on \mathbf{X} to select a few important features



PCA for classification can lead to problems...

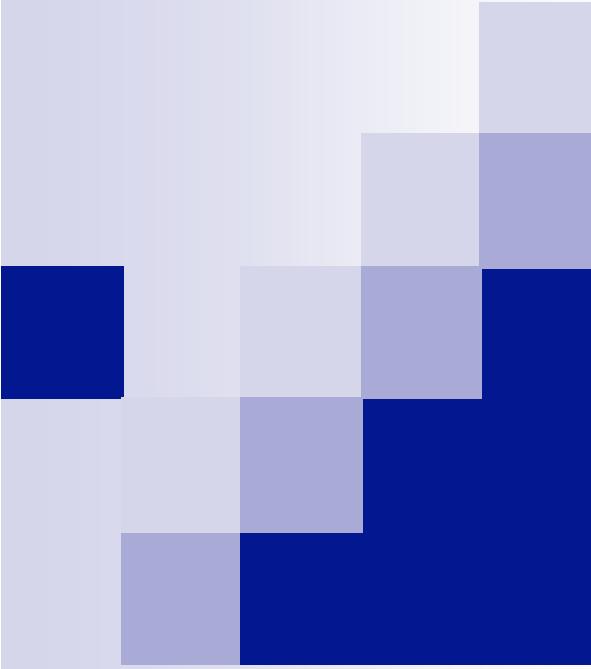
- Direction of maximum variation may be unrelated to “discriminative” directions:
- PCA often works very well, but sometimes must use more advanced methods
 - e.g., Fisher linear discriminant

What you need to know

- Dimensionality reduction
 - why and when it's important
- Simple feature selection
- Principal component analysis
 - minimizing reconstruction error
 - relationship to covariance matrix and eigenvectors
 - using SVD
 - problems with PCA

Announcements

- Homework 5:
 - Extension: Due Friday at 10:30am
 - Hand in to Monica, Wean 4619
- Project:
 - Poster session: Friday May 4th 2-5pm, NSH Atrium
 - please arrive a 15mins early to set up
 - Paper: Thursday May 10th by 2pm
 - electronic submission by email to instructors list
 - maximum of 8 pages, NIPS format
 - no late days allowed
- FCEs!!!!
 - Please, please, please, please, please, please give us your feedback, it helps us improve the class! ☺
 - <http://www.cmu.edu/fce>



Markov Decision Processes (MDPs)

Machine Learning – 10701/15781

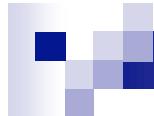
Carlos Guestrin

Carnegie Mellon University

April 25th, 2006

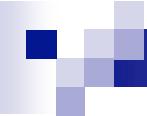
©2005-2007 Carlos Guestrin

Thus far this semester



- Regression:
- Classification:
- Density estimation:

Learning to act



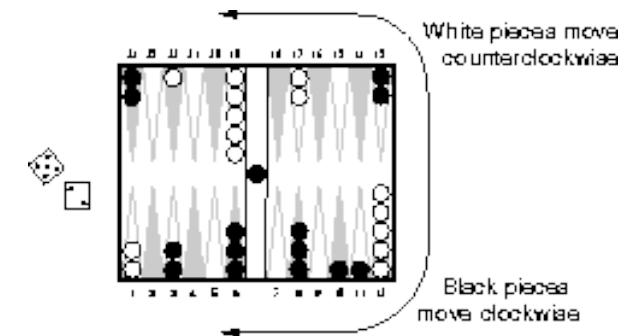
[Ng et al. '05]

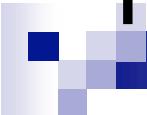
- Reinforcement learning
- An agent
 - Makes sensor observations
 - Must select action
 - Receives rewards
 - positive for “good” states
 - negative for “bad” states

Learning to play backgammon

[Tesauro '95]

- Combines reinforcement learning with neural networks
- Played 300,000 games against itself
- Achieved grandmaster level!

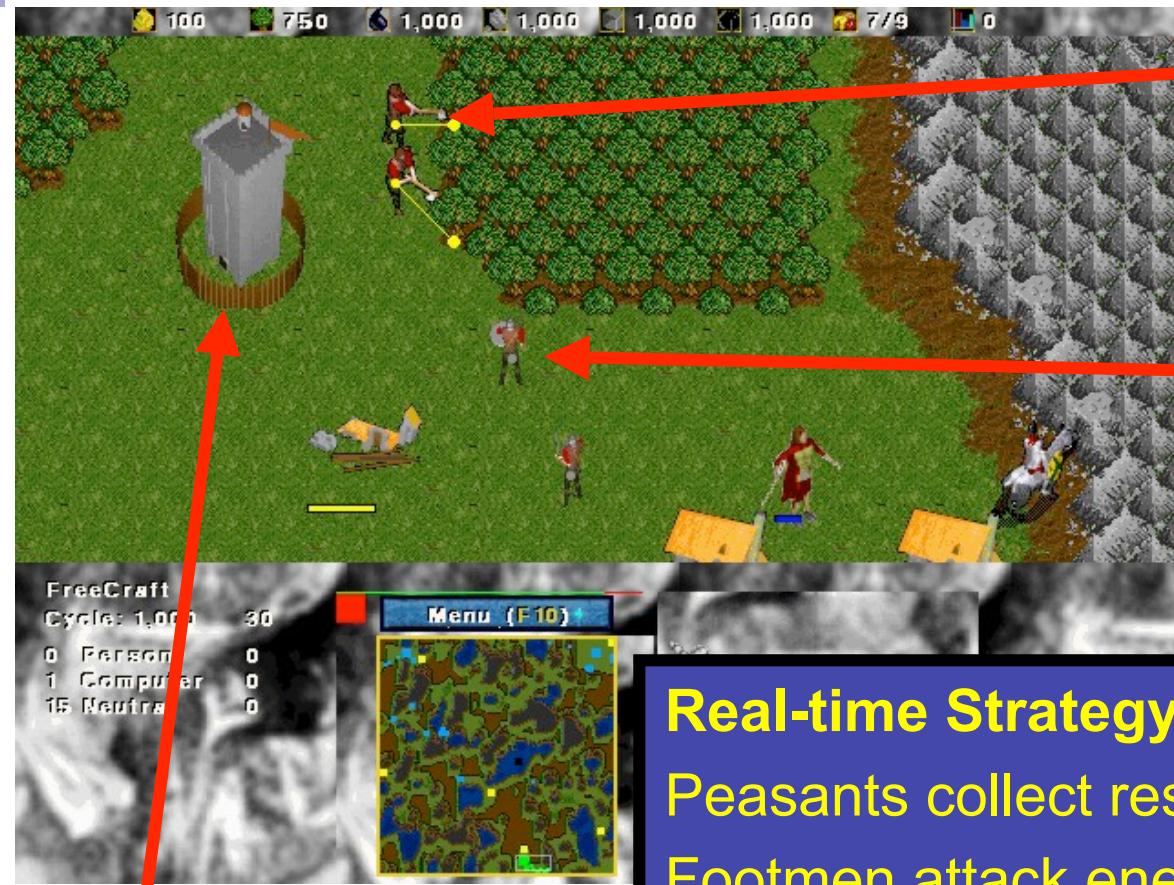




Roadmap to learning about reinforcement learning

- When we learned about Bayes nets:
 - First talked about formal framework:
 - representation
 - inference
 - Then learning for BNs
- For reinforcement learning:
 - Formal framework
 - Markov decision processes
 - Then learning

FreeCraft



building

peasant

footman

Real-time Strategy Game

Peasants collect resources and build
Footmen attack enemies
Buildings train peasants and footmen

States and actions

- State space:

- Joint state \mathbf{x} of entire system



- Action space:

- Joint action $\mathbf{a} = \{a_1, \dots, a_n\}$ for all agents

States change over time

- Like an HMM, state changes over time
- Next state depends on current state and action selected
 - e.g., action="build castle" likely to lead to a state where you have a castle
- Transition model:
 - Dynamics of the entire system $P(x'|x,a)$



Some states and actions are better than others

- Each state x is associated with a reward
 - positive reward for successful attack
 - negative for loss
- Reward function:
 - Total reward $R(x)$

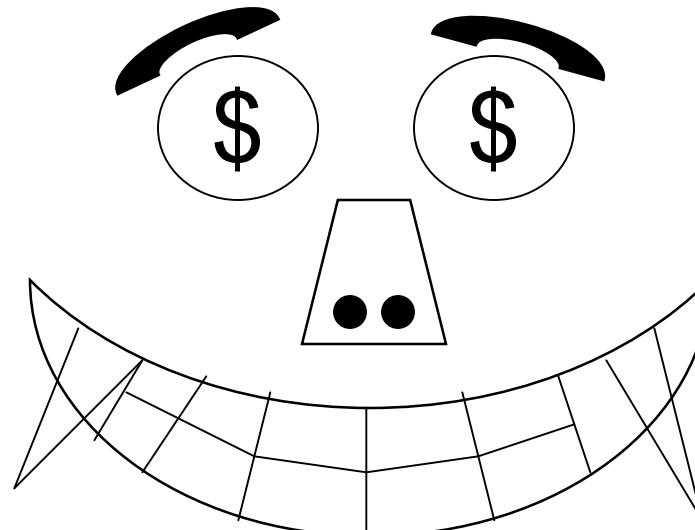


Discounted Rewards

An assistant professor gets paid, say, 20K per year.

How much, in total, will the A.P. earn in their life?

$$20 + 20 + 20 + 20 + 20 + \dots = \text{Infinity}$$



What's wrong with this argument?

Discounted Rewards



“A reward (payment) in the future is not worth quite as much as a reward now.”

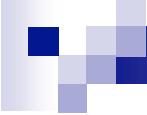
- Because of chance of obliteration
- Because of inflation

Example:

Being promised \$10,000 next year is worth only 90% as much as receiving \$10,000 right now.

Assuming payment n years in future is worth only $(0.9)^n$ of payment now, what is the AP's **Future Discounted Sum of Rewards** ?

Discount Factors



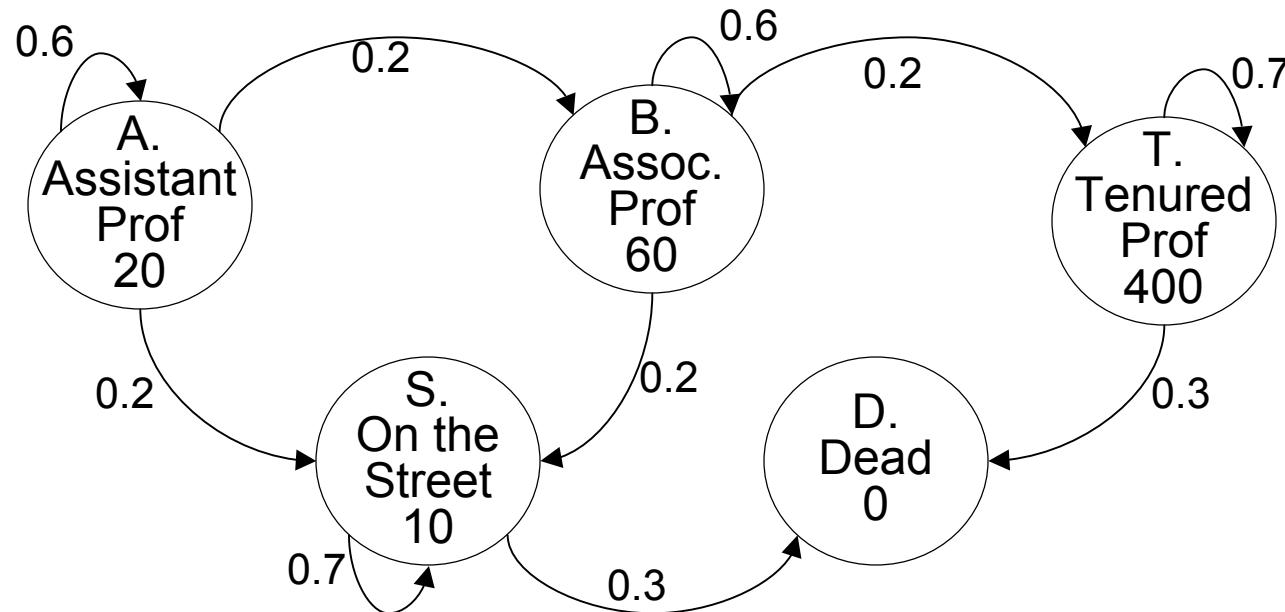
People in economics and probabilistic decision-making do this all the time.

The “Discounted sum of future rewards” using discount factor γ is

$$\begin{aligned} & (\text{reward now}) + \\ & \gamma (\text{reward in 1 time step}) + \\ & \gamma^2 (\text{reward in 2 time steps}) + \\ & \gamma^3 (\text{reward in 3 time steps}) + \\ & \vdots \\ & \quad : \quad (\text{infinite sum}) \end{aligned}$$

The Academic Life

Assume Discount Factor $\gamma = 0.9$



Define:

V_A = Expected discounted future rewards starting in state A

V_B = Expected discounted future rewards starting in state B

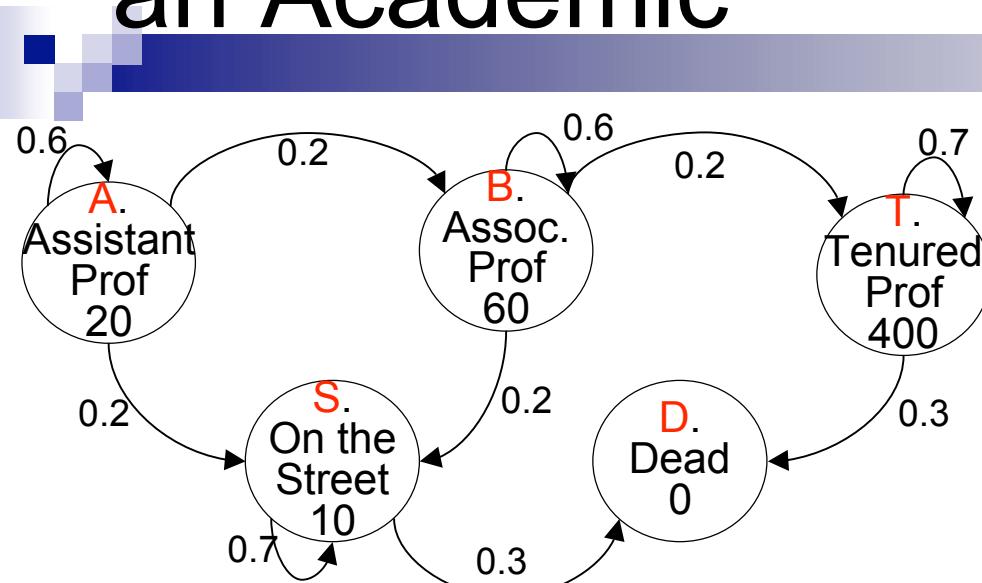
V_T = “ “ “ “ “ “ “ “ T

V_S = “ “ “ “ “ “ “ “ S

V_D = “ “ “ “ “ “ “ “ D

How do we compute V_A , V_B , V_T , V_S , V_D ?

Computing the Future Rewards of an Academic



Assume Discount Factor $\gamma = 0.9$

Joint Decision Space

Markov Decision Process (MDP) Representation:

- State space:
 - Joint state \mathbf{x} of entire system
- Action space:
 - Joint action $\mathbf{a} = \{a_1, \dots, a_n\}$ for all agents
- Reward function:
 - Total reward $R(\mathbf{x}, \mathbf{a})$
 - sometimes reward can depend on action
- Transition model:
 - Dynamics of the entire system $P(\mathbf{x}'|\mathbf{x}, \mathbf{a})$



Policy

Policy: $\pi(\mathbf{x}) = \mathbf{a}$

At state \mathbf{x} ,
action \mathbf{a} for all
agents



$\pi(\mathbf{x}_0)$ = both peasants get wood



$\pi(\mathbf{x}_1)$ = one peasant builds
barrack, other gets gold



$\pi(\mathbf{x}_2)$ = peasants get gold,
footmen attack

Value of Policy

Value: $V_\pi(x)$

Expected long-term reward starting from x

Start from x_0



$R(x_0)$

$\pi(x_0)$

$$V_\pi(x_0) = E_\pi[R(x_0) + \gamma R(x_1) + \gamma^2 R(x_2) + \gamma^3 R(x_3) + \gamma^4 R(x_4) + L]$$

Future rewards discounted by γ 2 [0,1)

$\pi(x_1)$

$R(x_1)$



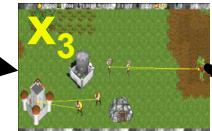
$R(x_1')$

$\pi(x_1'')$



$\pi(x_2)$

$R(x_2)$



$R(x_3)$

$\pi(x_3)$



$R(x_4)$

39

Computing the value of a policy


$$V_{\pi}(x_0) = E_{\pi}[R(x_0) + \gamma R(x_1) + \gamma^2 R(x_2) + \gamma^3 R(x_3) + \gamma^4 R(x_4) + L]$$

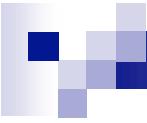
- Discounted value of a state:

- value of starting from x_0 and continuing with policy π from then on

$$\begin{aligned} V_{\pi}(x_0) &= E_{\pi}[R(x_0) + \gamma R(x_1) + \gamma^2 R(x_2) + \gamma^3 R(x_3) + \dots] \\ &= E_{\pi}\left[\sum_{t=0}^{\infty} \gamma^t R(x_t)\right] \end{aligned}$$

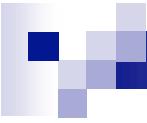
- A recursion!

Computing the value of a policy 1 – the matrix inversion approach


$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$

- Solve by simple matrix inversion:

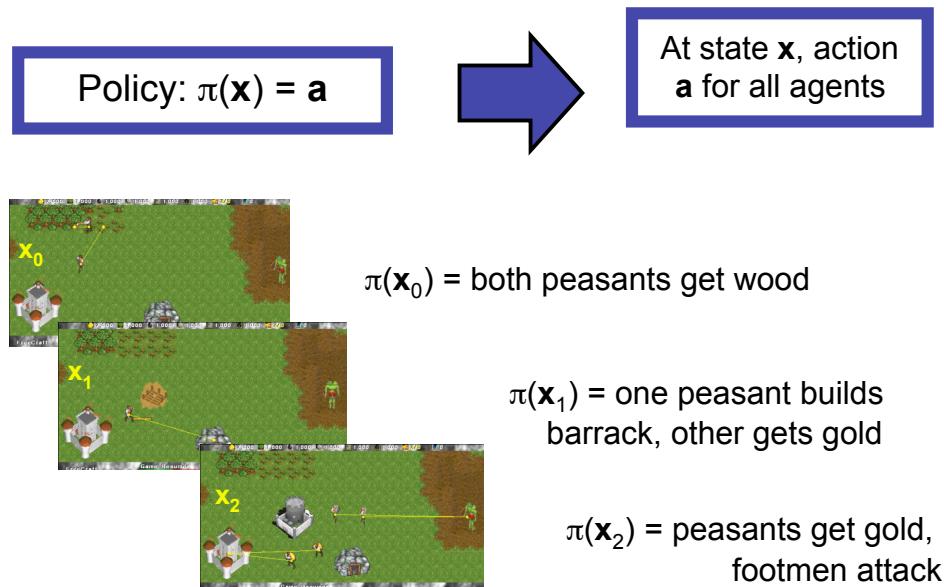
Computing the value of a policy 2 – iteratively


$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$

- If you have 1000,000 states, inverting a 1000,000x1000,000 matrix is hard!
- Can solve using a simple convergent iterative approach:
(a.k.a. dynamic programming)
 - Start with some guess V_0
 - Iteratively say:
 - $V_{t+1} = R + \gamma P_\pi V_t$
 - Stop when $\|V_{t+1} - V_t\|_1 \cdot \epsilon$
 - means that $\|V_\pi - V_{t+1}\|_1 \cdot \epsilon / (1 - \gamma)$

But we want to learn a Policy

- So far, told you how good a policy is...
- But how can we choose the best policy???
- Suppose there was only one time step:
 - world is about to end!!!
 - select action that maximizes reward!



Another recursion!

- Two time steps: address tradeoff
 - good reward now
 - better reward in the future

Unrolling the recursion

- Choose actions that lead to best value in the long run
 - Optimal value policy achieves optimal value V^*

$$V^*(x_0) = \max_{a_0} R(x_0, a_0) + \gamma E_{a_0} [\max_{a_1} R(x_1) + \gamma^2 E_{a_1} [\max_{a_2} R(x_2) + \dots]]$$

Bellman equation

- Evaluating policy π :

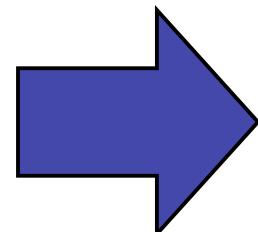
$$V_\pi(x) = R(x) + \gamma \sum_{x'} P(x' | x, a = \pi(x)) V_\pi(x')$$

- Computing the optimal value V^* - Bellman equation

$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Optimal Long-term Plan

Optimal value
function $V^*(\mathbf{x})$



Optimal Policy: $\pi^*(\mathbf{x})$

$$Q^*(\mathbf{x}, \mathbf{a}) = R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Optimal policy:

$$\pi^*(\mathbf{x}) = \arg \max_{\mathbf{a}} Q^*(\mathbf{x}, \mathbf{a})$$

Interesting fact – Unique value

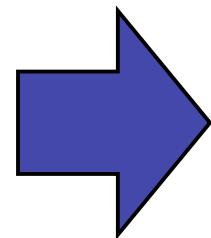

$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

- *Slightly surprising fact:* There is only one V^* that solves Bellman equation!
 - there may be many optimal policies that achieve V^*
- *Surprising fact:* optimal policies are good everywhere!!!

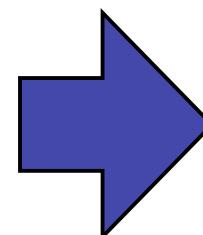
$$V_{\pi^*}(x) \geq V_\pi(x), \quad \forall x, \quad \forall \pi$$

Solving an MDP

Solve
Bellman
equation



Optimal
value $V^*(\mathbf{x})$



Optimal
policy $\pi^*(\mathbf{x})$

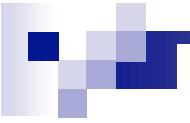
$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

Bellman equation is non-linear!!!

Many algorithms solve the Bellman equations:

- Policy iteration [Howard '60, Bellman '57]
- Value iteration [Bellman '57]
- Linear programming [Manne '60]
- ...

Value iteration (a.k.a. dynamic programming) – the simplest of all


$$V^*(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V^*(\mathbf{x}')$$

- Start with some guess V_0
- Iteratively say:

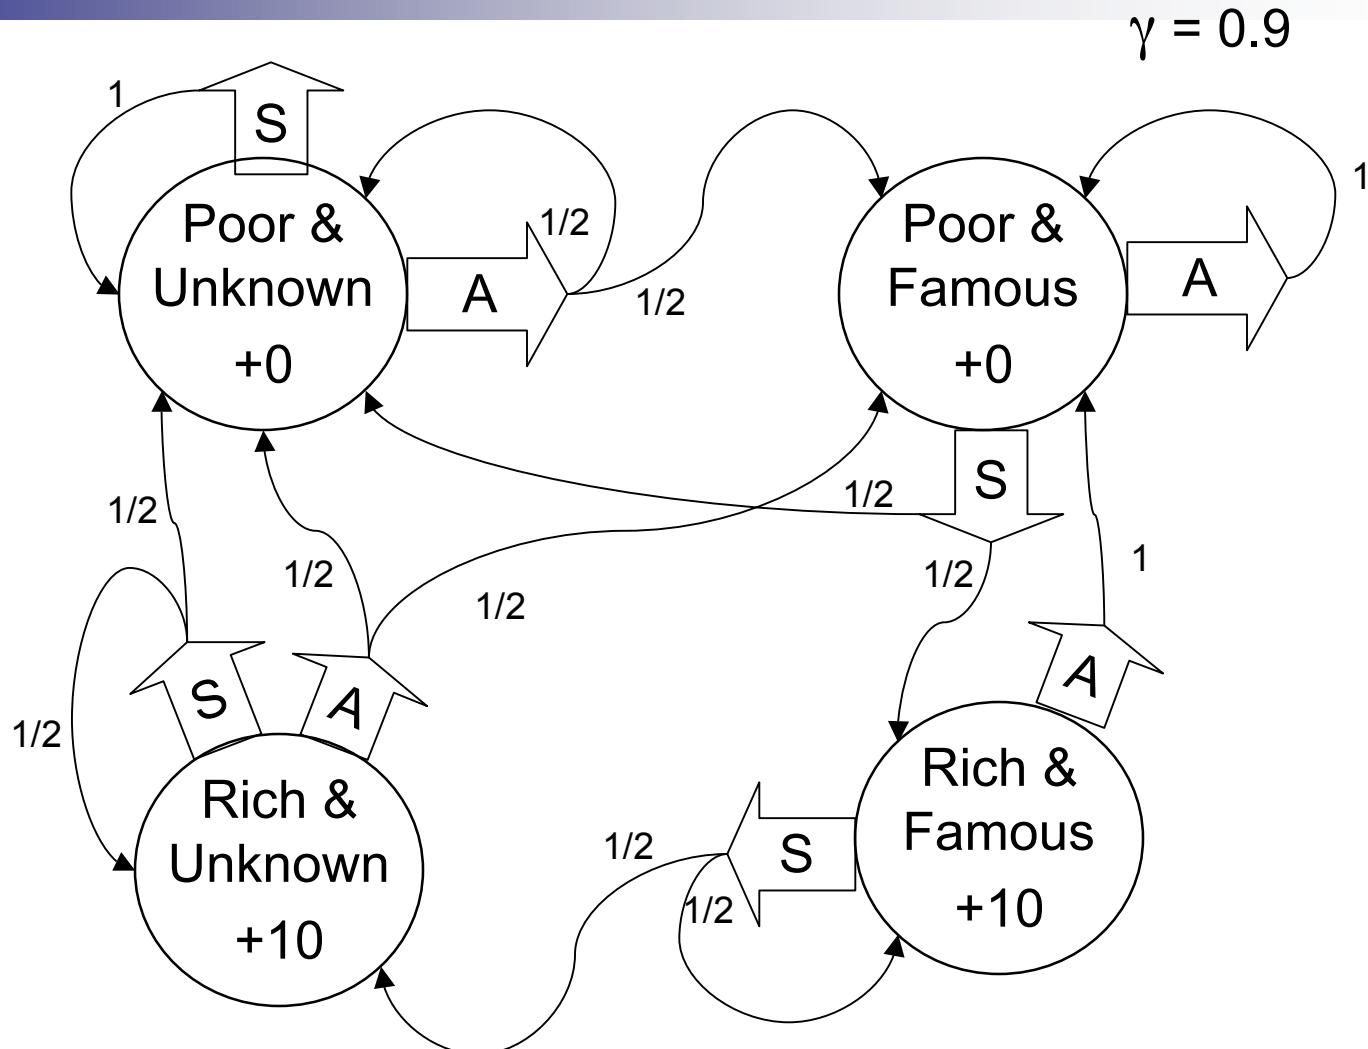
- $V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$

- Stop when $\|V_{t+1} - V_t\|_1 \leq \epsilon$
 - means that $\|V^* - V_{t+1}\|_1 \leq \epsilon/(1-\gamma)$

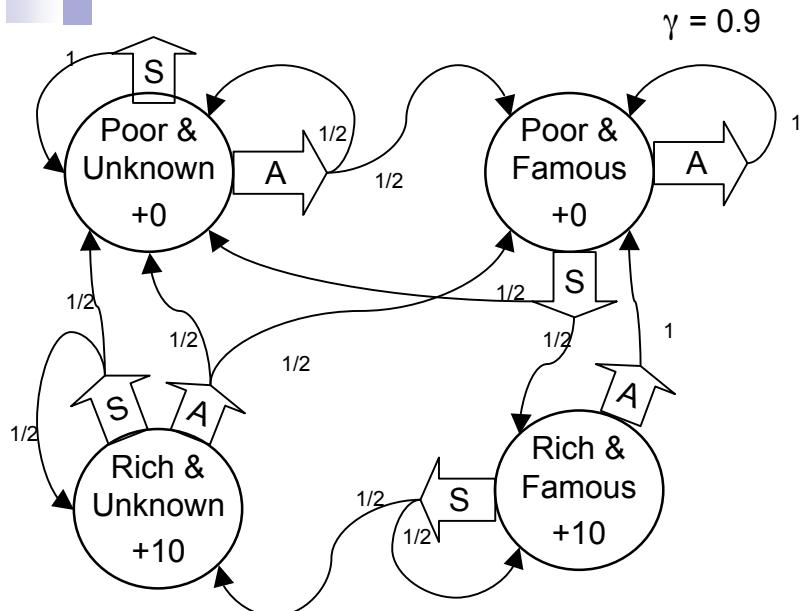
A simple example

You run a startup company.

In every state you must choose between Saving money or Advertising.



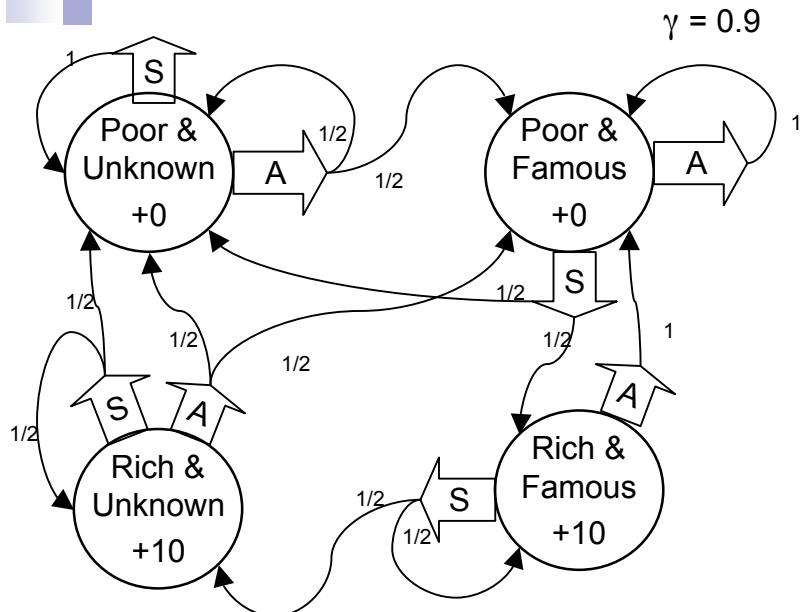
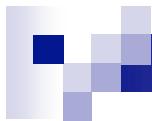
Let's compute $V_t(x)$ for our example



t	$V_t(\text{PU})$	$V_t(\text{PF})$	$V_t(\text{RU})$	$V_t(\text{RF})$
1				
2				
3				
4				
5				
6				

$$V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

Let's compute $V_t(x)$ for our example



t	$V_t(\text{PU})$	$V_t(\text{PF})$	$V_t(\text{RU})$	$V_t(\text{RF})$
1	0	0	10	10
2	0	4.5	14.5	19
3	2.03	6.53	25.08	18.55
4	3.852	12.20	29.63	19.26
5	7.22	15.07	32.00	20.40
6	10.03	17.65	33.58	22.43

$$V_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

Policy iteration – Another approach for computing π^*

- Start with some guess for a policy π_0
- Iteratively say:

- evaluate policy:

$$V_t(\mathbf{x}) = R(\mathbf{x}, \mathbf{a} = \pi_t(\mathbf{x})) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a} = \pi_t(\mathbf{x})) V_t(\mathbf{x}')$$

- improve policy:

$$\pi_{t+1}(\mathbf{x}) = \max_{\mathbf{a}} R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V_t(\mathbf{x}')$$

- Stop when
 - policy stops changing
 - usually happens in about 10 iterations
 - or $\|V_{t+1} - V_t\|_1 \cdot \epsilon$
 - means that $\|V^* - V_{t+1}\|_1 \cdot \epsilon / (1 - \gamma)$

Policy Iteration & Value Iteration: Which is best ???

It depends.

Lots of actions? Choose **Policy Iteration**

Already got a fair policy? **Policy Iteration**

Few actions, acyclic? **Value Iteration**

Best of Both Worlds:

Modified Policy Iteration [Puterman]

...a simple mix of value iteration and policy iteration

3rd Approach

Linear Programming

LP Solution to MDP

[Manne '60]

Value computed by linear programming:

$$\text{minimize: } \sum_{\mathbf{x}} V(\mathbf{x})$$

$$\text{subject to : } \begin{cases} V(\mathbf{x}) \geq R(\mathbf{x}, \mathbf{a}) + \gamma \sum_{\mathbf{x}'} P(\mathbf{x}' | \mathbf{x}, \mathbf{a}) V(\mathbf{x}') \\ \forall \mathbf{x}, \mathbf{a} \end{cases}$$

- One variable $V(\mathbf{x})$ for each state
- One constraint for each state \mathbf{x} and action \mathbf{a}
- **Polynomial time solution**

What you need to know

- What's a Markov decision process
 - state, actions, transitions, rewards
 - a policy
 - value function for a policy
 - computing V_π
- Optimal value function and optimal policy
 - Bellman equation
- Solving Bellman equation
 - with value iteration, policy iteration and linear programming

Acknowledgment

- This lecture contains some material from Andrew Moore's excellent collection of ML tutorials:
 - <http://www.cs.cmu.edu/~awm/tutorials>